Deep Convolutional Generative Adversarial Network

Alireza Javaheri

OVERVIEW

Objectives

Architecture

Datasets

Training

Results

OBJECTIVES

Propose and evaluate a set of constraints on the architectural topology of Convolutional GANs that make them stable to train in most settings. We name this class of architectures Deep Convolutional GANs (DCGAN)

Use the trained discriminators for image classification tasks, showing competitive performance with other unsupervised algorithms

We visualize the filters learnt by GANs and empirically show that specific filters have learned to draw specific objects

ARCHITECTURE

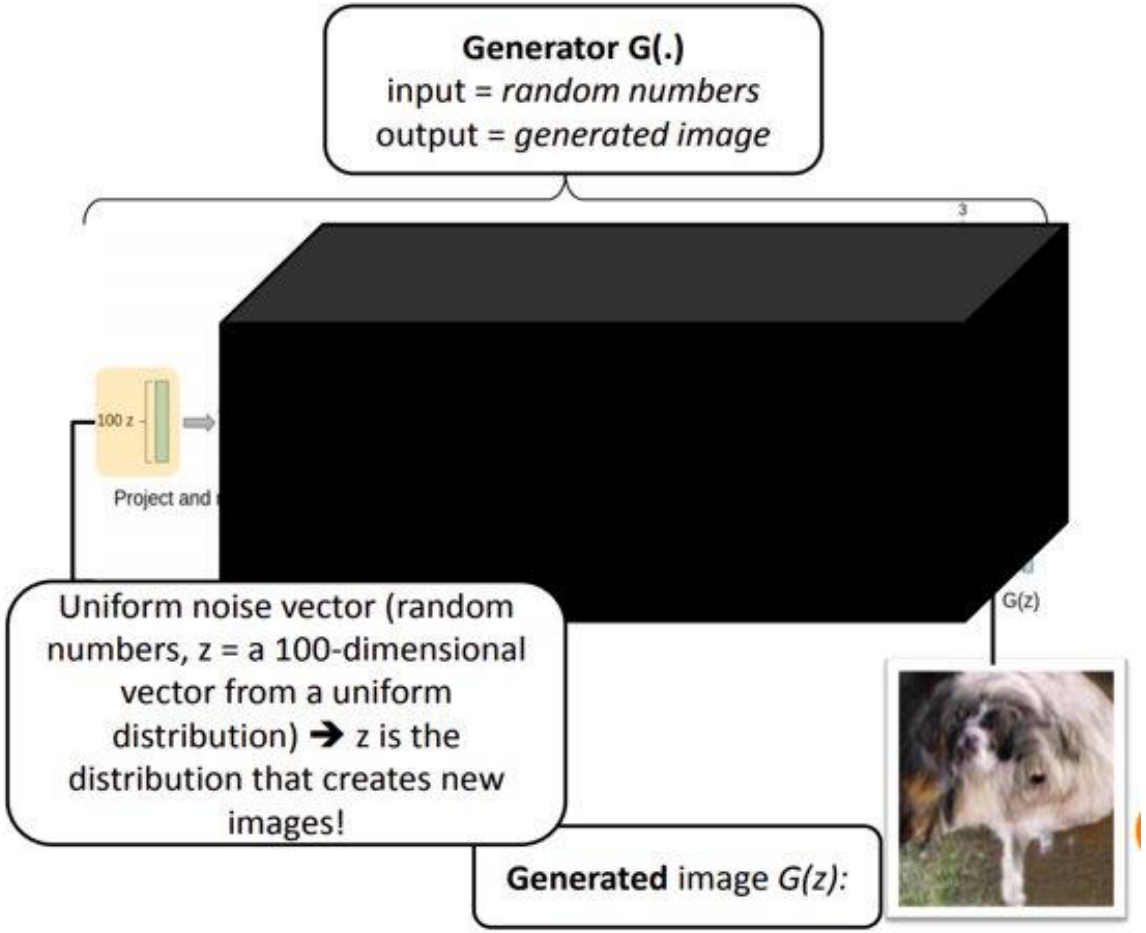

ARCHITECTURE

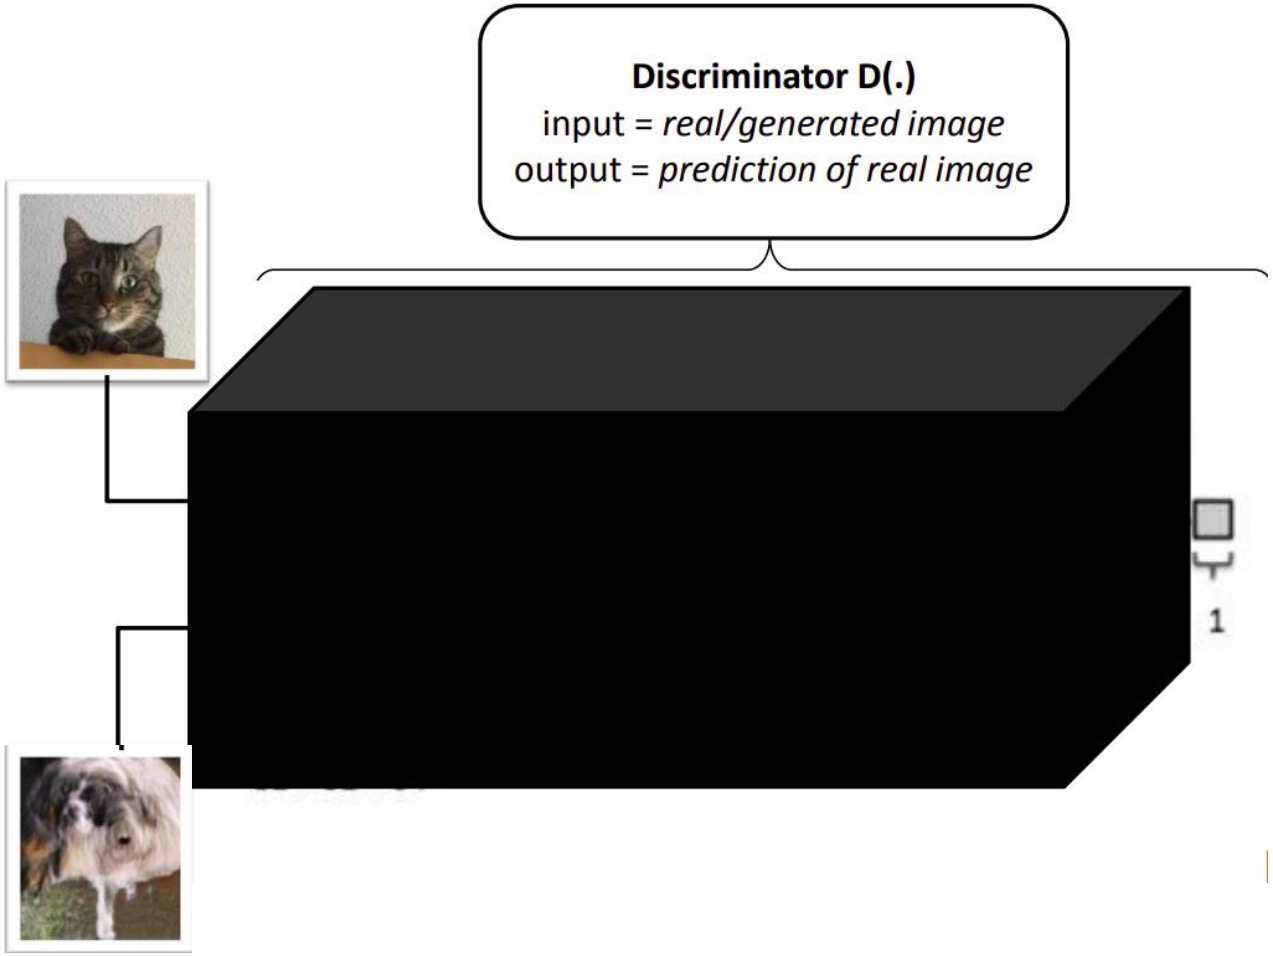

ARCHITECTURE

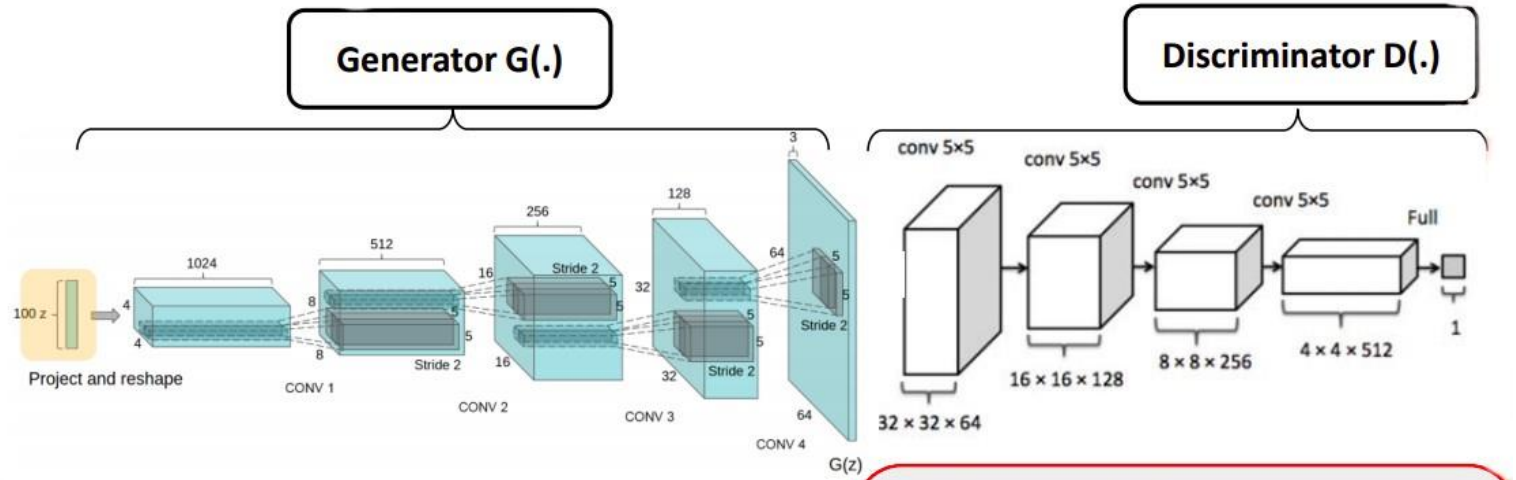

Generator Goal: Fool *D(G(z))*
i.e., generate an image *G(z)* such that *D(G(z))* is wrong, i.e., *D(G(z)) = 1*.

Discriminator Goal: discriminate between **real** and **generated** images *i.e., D(x)=1,* where *x* is a **real** image *D(G(z))=0,* where *G(z)* is a **generated** image.

- **Conflicting** goals.
- Both goals are **unsupervised.**
- Optimal when D(.)=0.5 (i.e., cannot tell the difference between real and generated images) and G(z)=learns the training images distribution.

ARCHITECTURE

Generator

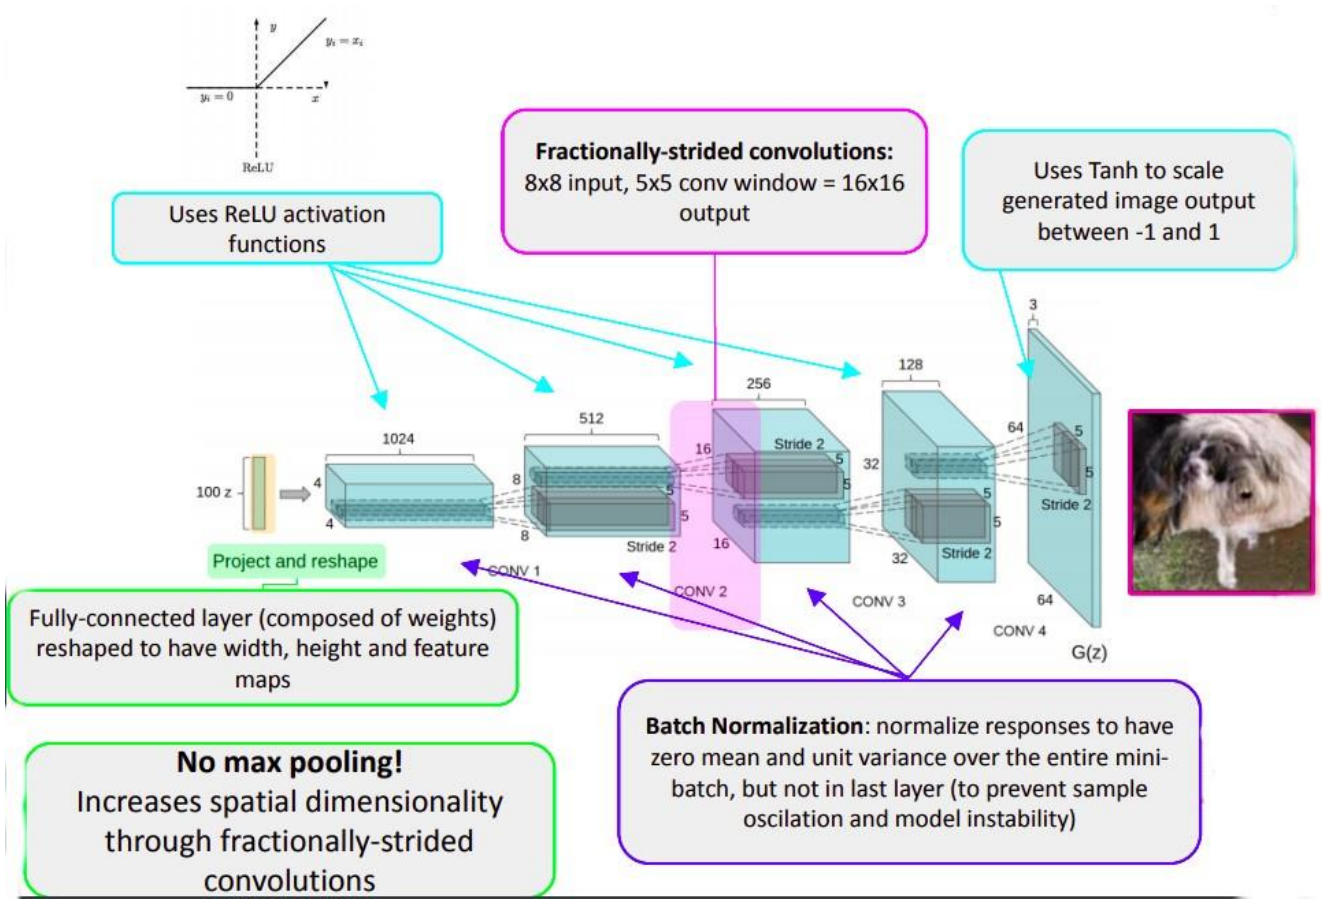

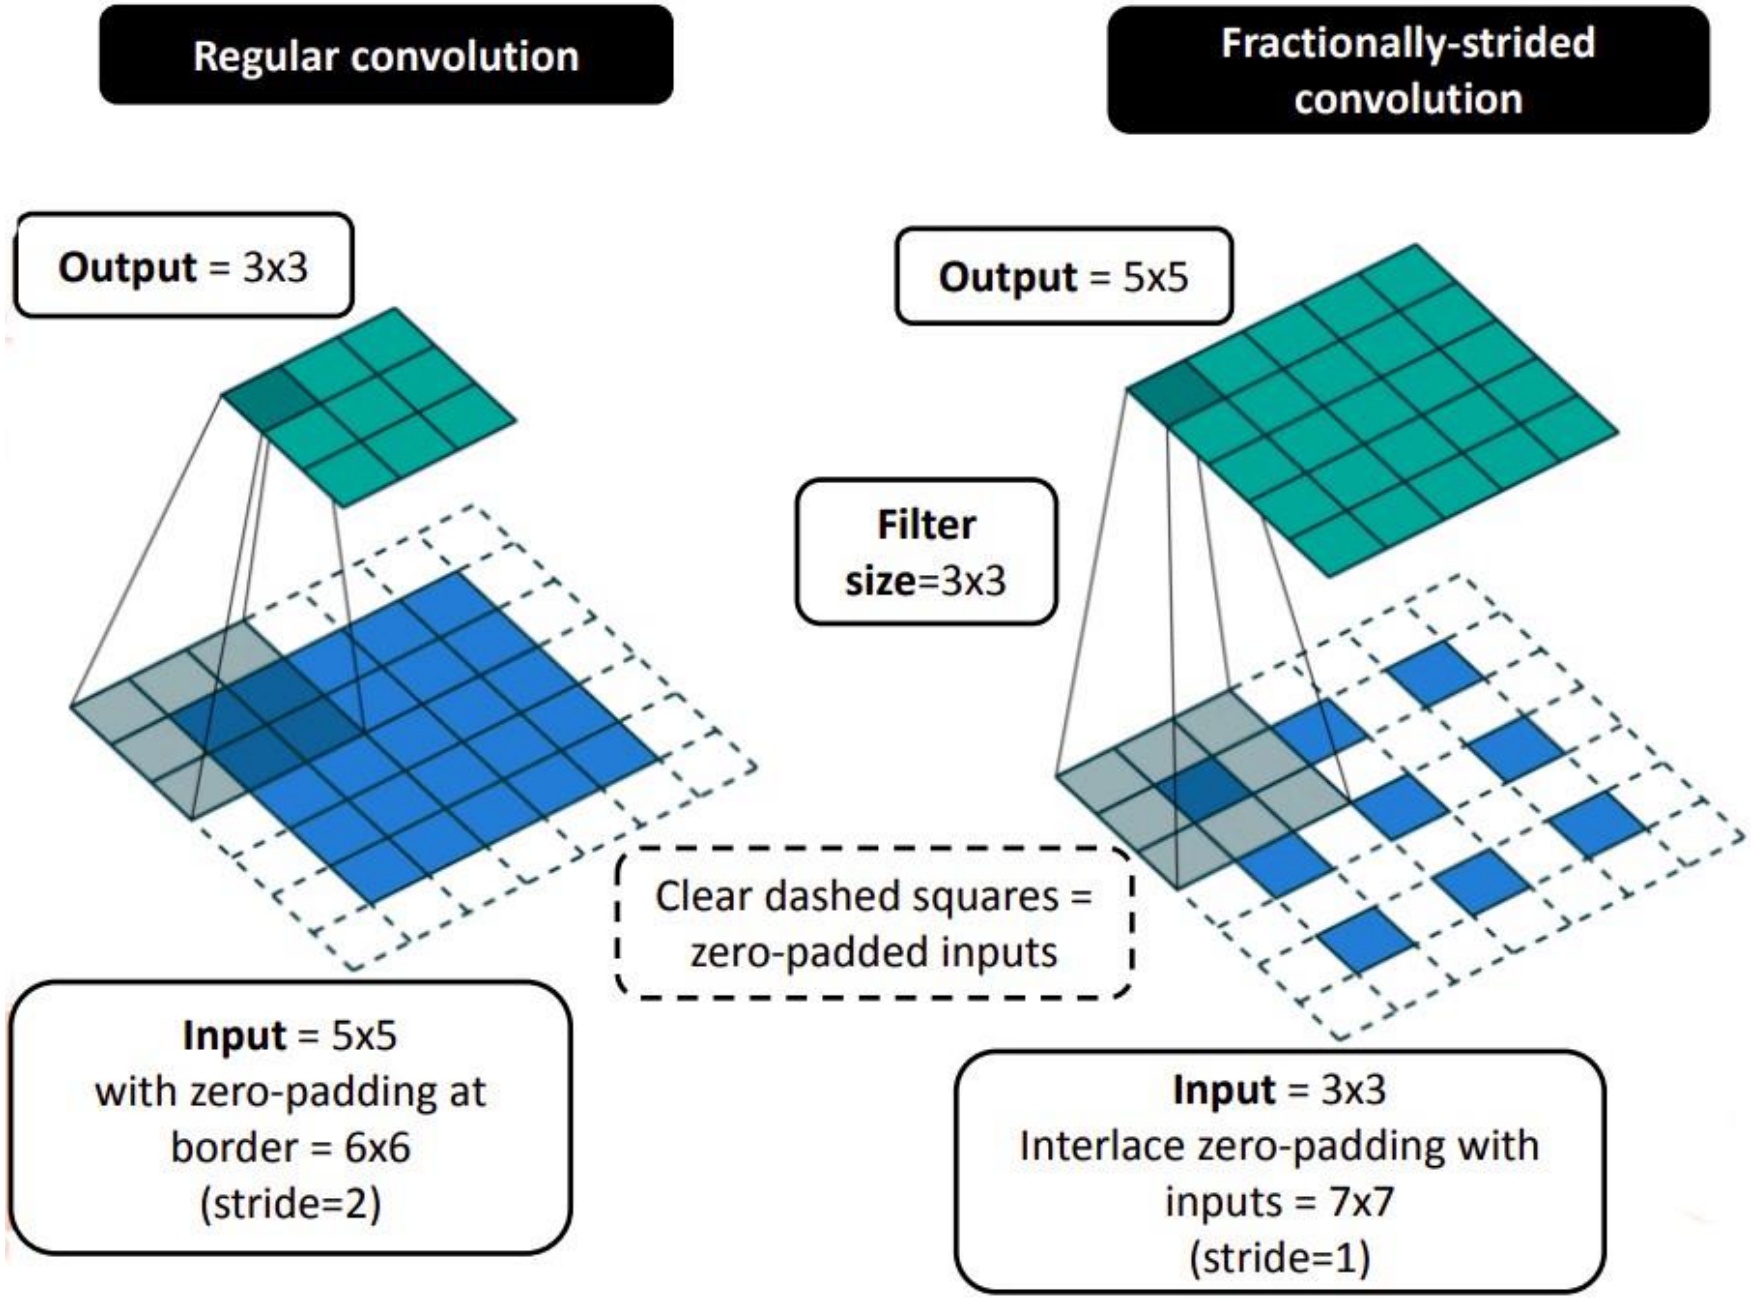
Regular convolution
Fractionally-strided convolution
Output = 3x3
Output = 5x5
Filter size=3x3
Clear dashed squares = zero-padded inputs
Input = 5x5 with zero-padding at border = 6x6 (stride=2)
Input = 3x3 Interlace zero-padding with inputs = 7x7 (stride=1)

ARCHITECTURE

Discriminator

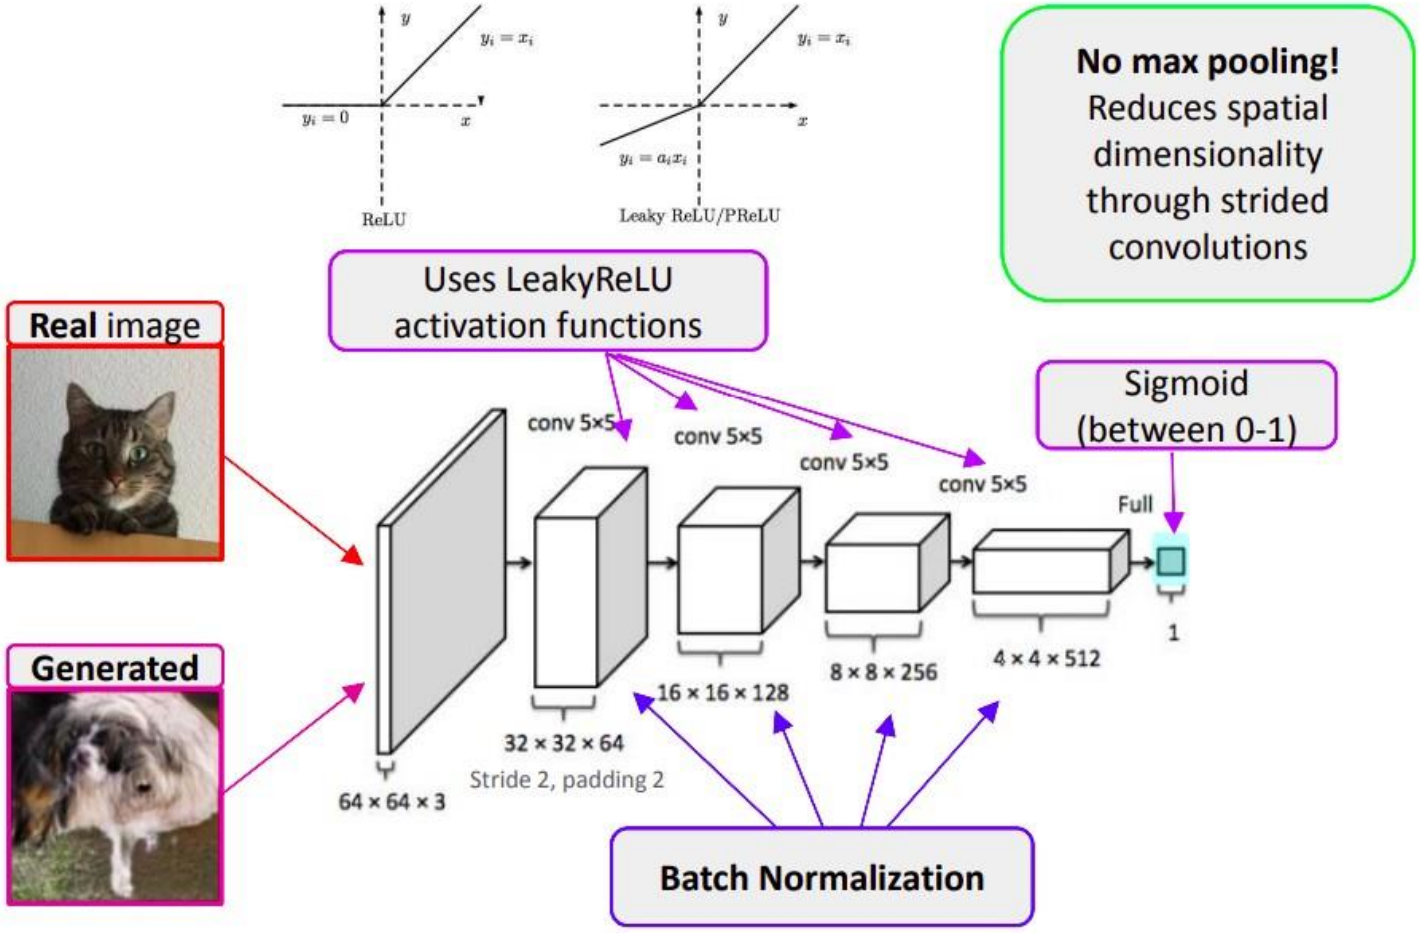

ARCHITECTURE GUIDELINES FOR STABLE DEEP CONVOLUTIONAL GANS

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Use ReLU activation in generator
- Use LeakyReLU activation in the discriminator.

Datasets

- Large-scale Scene Understanding (LSUN)
- Faces
- ImageNet-1K

TRAINING

- Pre-processing: scale images between -1 and 1 (tanh range).
- Minibatch SGD (m = 128).
- Weight init.: zero-centered normal distribution (std. dev. = 0.02).
- Leaky ReLU slope = 0.2.
- Adam optimizer with tuned hyperparameters to accelerate training.
- Learning rate = 0.0002.
- Momentum term $\beta 1 = 0.5$ to stabilize training

Results

Results

INTERPOLATION BETWEEN A SERIES OF RANDOM POINTS IN Z

Results

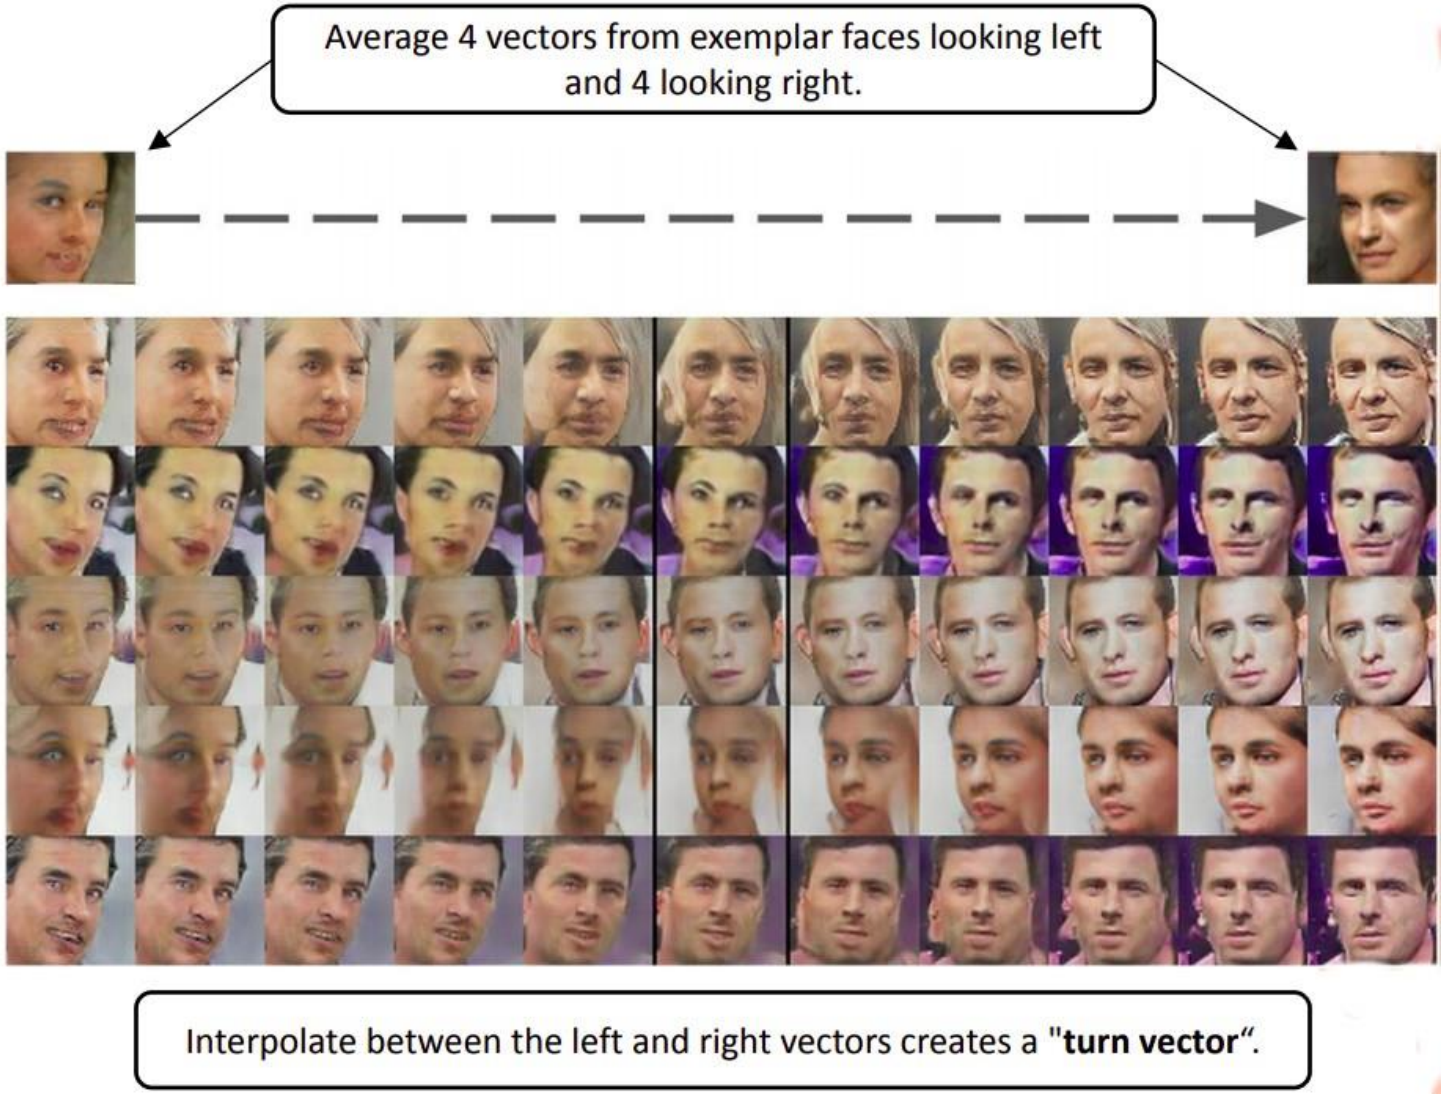
Average 4 vectors from exemplar faces looking left and 4 looking right.
Interpolate between the left and right vectors creates a "**turn vector**".

Results

(Bottom) Samples generated after dropping out "window" concept.
Some windows are removed or transformed.
The overall scene stays the same, indicating the generator has separated objects (windows) from the scene.

Results

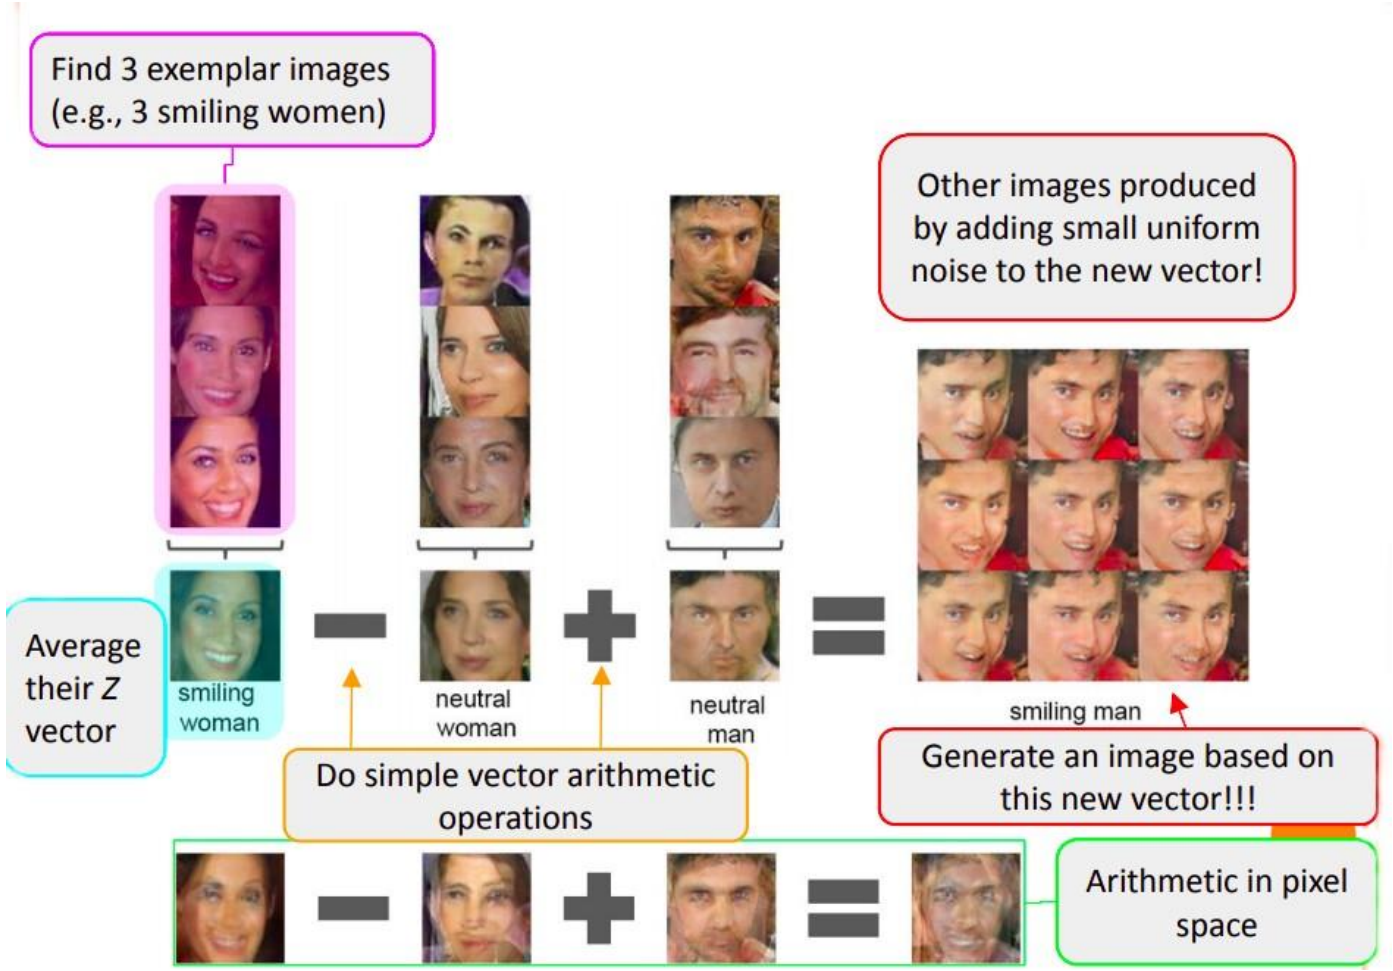
Find 3 exemplar images
(e.g., 3 smiling women)
Other images produced by adding small uniform noise to the new vector!
Average their Z vector
smiling woman
neutral woman
neutral man
smiling man
Do simple vector arithmetic operations
Generate an image based on this new vector!!!
Arithmetic in pixel space